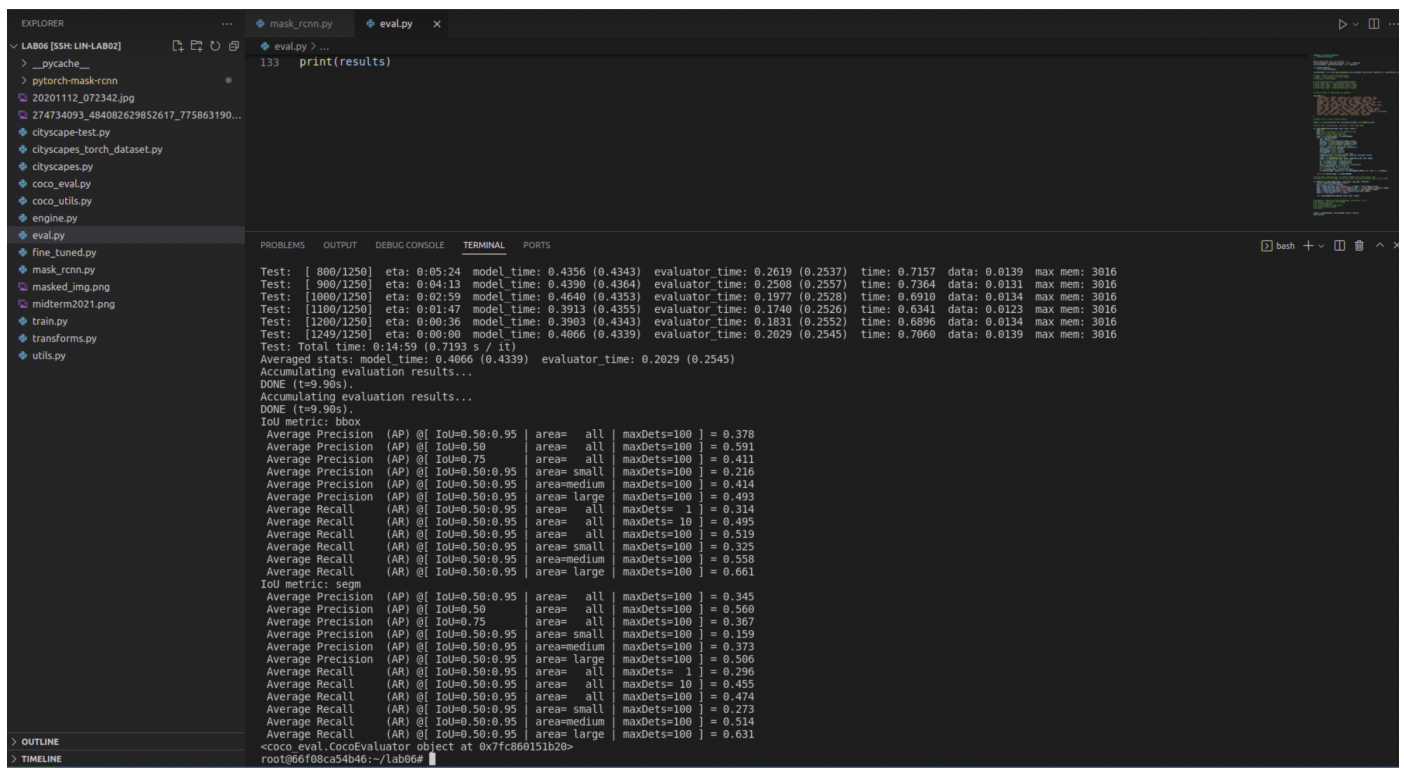


Lab06 Report

st122314

In this report I learnt about Masked-RCNN using COCO dataset. Moreover, I also tested with Cityscapes data with Pre-train model and then learnt how to make fine-tuning.

The following figure is the report after running COCO eval.py.



In []:

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

In []:

```
cd /content/drive/MyDrive/RTML/Midsem_prep/data
/content/drive/MyDrive/RTML/Midsem_prep/data
```

CityScape data download

In []:

```
!wget --keep-session-cookies --save-cookies=cookies.txt --post-data 'username=winwin'
--2022-03-03 10:37:54-- https://www.cityscapes-dataset.com/login/ (ht
tps://www.cityscapes-dataset.com/login/)
Resolving www.cityscapes-dataset.com (www.cityscapes-dataset.com)... 1
39.19.217.8
Connecting to www.cityscapes-dataset.com (www.cityscapes-dataset.com)|
139.19.217.8|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://www.cityscapes-dataset.com/downloads/ (https://www.c
ityscapes-dataset.com/downloads/) [following]
--2022-03-03 10:37:55-- https://www.cityscapes-dataset.com/downloads/
(https://www.cityscapes-dataset.com/downloads/)
Reusing existing connection to www.cityscapes-dataset.com:443.
HTTP request sent, awaiting response... 200 OK
Length: unspecified [text/html]
Saving to: 'index.html.1'

index.html.1          [ <=>          ] 47.44K  --.-KB/s    in
0.1s

2022-03-03 10:37:55 (393 KB/s) - 'index.html.1' saved [48580]
```

In []:

```
!wget --load-cookies cookies.txt --content-disposition https://www.cityscapes-dataset
--2022-03-03 10:37:59-- https://www.cityscapes-dataset.com/file-handl
ing/?packageID=1 (https://www.cityscapes-dataset.com/file-handling/?pa
ckageID=1)
Resolving www.cityscapes-dataset.com (www.cityscapes-dataset.com)... 1
39.19.217.8
Connecting to www.cityscapes-dataset.com (www.cityscapes-dataset.com)|
139.19.217.8|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 252567705 (241M) [application/octet-stream]
Saving to: 'gtFine_trainvaltest.zip'

gtFine_trainvaltest 100%[=====>] 240.87M  24.9MB/s    in
11s

2022-03-03 10:38:11 (21.2 MB/s) - 'gtFine_trainvaltest.zip' saved [252
567705/252567705]
```

In []:

```
!wget --load-cookies cookies.txt --content-disposition https://www.cityscapes-dataset.com/file-handling/?packageID=3 (https://www.cityscapes-dataset.com/file-handling/?packageID=3)
Resolving www.cityscapes-dataset.com (www.cityscapes-dataset.com)... 139.19.217.8
Connecting to www.cityscapes-dataset.com (www.cityscapes-dataset.com)|139.19.217.8|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 11592327197 (11G) [application/octet-stream]
Saving to: 'leftImg8bit_trainvaltest.zip'

leftImg8bit_trainva 100%[=====>] 10.80G 22.2MB/s in
7m 24s

2022-03-03 10:45:55 (24.9 MB/s) - 'leftImg8bit_trainvaltest.zip' saved
[11592327197/11592327197]
```

In []:

```
ls

AlexNet.ipynb          GAN.ipynb             List.gdoc
cityscape_mask_rcnn.ipynb  GoogLeNet.ipynb      Masked-R-C
NN.ipynb               index.html             'Resnet&SEn
et.ipynb'              index.html.1          YOLO.ipynb
cookies.txt            leftImg8bit_trainvaltest.zip
data/
```

In []:

```
cd data/

/content/drive/MyDrive/RTML/Midsem_prep/data
```

In []:

!unzip gtFine_trainvaltest.zip

Streaming output truncated to the last 5000 lines.

```

inflating: gtFine/test/berlin/berlin_000117_000019_gtFine_color.png
inflating: gtFine/test/berlin/berlin_000114_000019_gtFine_color.png
inflating: gtFine/test/berlin/berlin_000434_000019_gtFine_labelIds.p
ng
inflating: gtFine/test/berlin/berlin_000420_000019_gtFine_color.png
inflating: gtFine/test/berlin/berlin_000483_000019_gtFine_instanceId
s.png
inflating: gtFine/test/berlin/berlin_000420_000019_gtFine_instanceId
s.png
inflating: gtFine/test/berlin/berlin_000254_000019_gtFine_color.png
inflating: gtFine/test/berlin/berlin_000490_000019_gtFine_color.png
inflating: gtFine/test/berlin/berlin_000448_000019_gtFine_polygons.j
son
inflating: gtFine/test/berlin/berlin_000099_000019_gtFine_labelIds.p
ng
inflating: gtFine/test/berlin/berlin_000068_000019_gtFine_instanceId
s.png
inflating: gtFine/test/berlin/berlin_000288_000019_gtFine_instanceId

```

In []:

!unzip leftImg8bit_trainvaltest.zip

Streaming output truncated to the last 5000 lines.

```

extracting: leftImg8bit/train/jena/jena_000074_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000040_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000020_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000030_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000005_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000059_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000100_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000034_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000089_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000104_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000107_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000080_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000006_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000082_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000044_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000026_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000115_000019_leftImg8bit.png
extracting: leftImg8bit/train/jena/jena_000023_000019_leftImg8bit.png

```

In []:

Cityscape run on mask rcnn

In []:

```
import torch
import torchvision
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor
from torchvision.datasets import Cityscapes

print('Loading pretrained model...')
model = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained=True).cuda()
model.eval()
```

Loading pretrained model...

In []:

```

dom
py as np
import Image
matplotlib.pyplot as plt
ch
.utils.data import Dataset, DataLoader
vision import transforms, utils
chvision.transforms.functional as TF
ctions import namedtuple

scapesDataset(Dataset):

init__(self, root, split='train', mode='fine', augment=False):

lf.root = os.path.expanduser(root)
lf.mode = 'gtFine' if mode == 'fine' else 'gtCoarse'
lf.images_dir = os.path.join(self.root, 'leftImg8bit', split)
lf.targets_dir = os.path.join(self.root, self.mode, split)
int(self.images_dir)
int(self.targets_dir)
lf.split = split
lf.augment = augment
lf.images = []
lf.targets = []
lf.mapping = {
    0: 0, # unlabeled
    1: 0, # ego vehicle
    2: 0, # rect border
    3: 0, # out of roi
    4: 0, # static
    5: 0, # dynamic
    6: 0, # ground
    7: 1, # road
    8: 0, # sidewalk
    9: 0, # parking
    10: 0, # rail track
    11: 0, # building
    12: 0, # wall
    13: 0, # fence
    14: 0, # guard rail
    15: 0, # bridge
    16: 0, # tunnel
    17: 0, # pole
    18: 0, # polegroup
    19: 0, # traffic light
    20: 0, # traffic sign
    21: 0, # vegetation
    22: 0, # terrain
    23: 2, # sky
    24: 0, # person
    25: 0, # rider
    26: 3, # car
    27: 0, # truck
    28: 0, # bus
    29: 0, # caravan
    30: 0, # trailer
    31: 0, # train

```

```

32: 0, # motorcycle
33: 0, # bicycle
-1: 0 # licenseplate

lf.mappingrgb = {
    0: (255, 0, 0), # unlabeled
    1: (255, 0, 0), # ego vehicle
    2: (255, 0, 0), # rect border
    3: (255, 0, 0), # out of roi
    4: (255, 0, 0), # static
    5: (255, 0, 0), # dynamic
    6: (255, 0, 0), # ground
    7: (0, 255, 0), # road
    8: (255, 0, 0), # sidewalk
    9: (255, 0, 0), # parking
    10: (255, 0, 0), # rail track
    11: (255, 0, 0), # building
    12: (255, 0, 0), # wall
    13: (255, 0, 0), # fence
    14: (255, 0, 0), # guard rail
    15: (255, 0, 0), # bridge
    16: (255, 0, 0), # tunnel
    17: (255, 0, 0), # pole
    18: (255, 0, 0), # polegroup
    19: (255, 0, 0), # traffic light
    20: (255, 0, 0), # traffic sign
    21: (255, 0, 0), # vegetation
    22: (255, 0, 0), # terrain
    23: (0, 0, 255), # sky
    24: (255, 0, 0), # person
    25: (255, 0, 0), # rider
    26: (255, 255, 0), # car
    27: (255, 0, 0), # truck
    28: (255, 0, 0), # bus
    29: (255, 0, 0), # caravan
    30: (255, 0, 0), # trailer
    31: (255, 0, 0), # train
    32: (255, 0, 0), # motorcycle
    33: (255, 0, 0), # bicycle
    -1: (255, 0, 0) # licenseplate

```

Ensure that this matches the above mapping!#!@#!@#

For example 4 classes, means we should map to the ids=(0,1,2,3)

This is used to specify how many outputs the network should product...

```
lf.num_classes = 4
```

```
=====
Check that inputs are valid
=====
```

```

mode not in ['fine', 'coarse']:
    raise ValueError('Invalid mode! Please use mode="fine" or mode="coarse"')
mode == 'fine' and split not in ['train', 'test', 'val']:
    raise ValueError('Invalid split for mode "fine"! Please use split="train", split="
if mode == 'coarse' and split not in ['train', 'train_extra', 'val']:
    raise ValueError('Invalid split for mode "coarse"! Please use split="train", split=
not os.path.isdir(self.images_dir) or not os.path.isdir(self.targets_dir):
    raise RuntimeError('Dataset not found or incomplete. Please make sure all required
        ' specified "split" and "mode" are inside the "root" directory'
=====

```

Read in the paths to all images

```
=====
r city in os.listdir(self.images_dir):
    img_dir = os.path.join(self.images_dir, city)
    target_dir = os.path.join(self.targets_dir, city)
    for file_name in os.listdir(img_dir):
        self.images.append(os.path.join(img_dir, file_name))
        target_name = '{}_{}'.format(file_name.split('_leftImg8bit')[0], '{}_labelIds.'.format(file_name.split('_leftImg8bit')[0]))
        # target_name = '{}_{}'.format(file_name.split('_leftImg8bit')[0], '{}_color.png'.format(file_name.split('_leftImg8bit')[0]))
        self.targets.append(os.path.join(target_dir, target_name))

repr__(self):
t_str = 'Dataset ' + self.__class__.__name__ + '\n'
t_str += '    Number of images: {}\n'.format(self.__len__())
t_str += '    Split: {}\n'.format(self.split)
t_str += '    Mode: {}\n'.format(self.mode)
t_str += '    Augment: {}\n'.format(self.augment)
t_str += '    Root Location: {}\n'.format(self.root)
turn fmt_str

len__(self):
turn len(self.images)

sk_to_class(self, mask):
'
ven the cityscapes dataset, this maps to a 0..classes numbers.
is is because we are using a subset of all masks, so we have this "mapping" function
is mapping function is used to map all the standard ids into the smaller subset.
'
skimg = torch.zeros((mask.size()[0], mask.size()[1]), dtype=torch.uint8)
r k in self.mapping:
    maskimg[mask == k] = self.mapping[k]
turn maskimg

sk_to_rgb(self, mask):
'
ven the Cityscapes mask file, this converts the ids into rgb colors.
is is needed as we are interested in a sub-set of labels, thus can't just use the
andard color output provided by the dataset.
'
bimg = torch.zeros((3, mask.size()[0], mask.size()[1]), dtype=torch.uint8)
r k in self.mappingrgb:
    rgbimg[0][mask == k] = self.mappingrgb[k][0]
    rgbimg[1][mask == k] = self.mappingrgb[k][1]
    rgbimg[2][mask == k] = self.mappingrgb[k][2]
turn rgbimg

ass_to_rgb(self, mask):
'
is function maps the classification index ids into the rgb.
r example after the argmax from the network, you want to find what class
given pixel belongs too. This does that but just changes the color
that we can compare it directly to the rgb groundtruth label.
'
sk2class = dict((v, k) for k, v in self.mapping.items())
bimg = torch.zeros((3, mask.size()[0], mask.size()[1]), dtype=torch.uint8)
r k in mask2class:
    rgbimg[0][mask == k] = self.mappingrgb[mask2class[k]][0]
    rgbimg[1][mask == k] = self.mappingrgb[mask2class[k]][1]
    rgbimg[2][mask == k] = self.mappingrgb[mask2class[k]][2]
turn rgbimg
```



```

getitem__(self, index):

    first load the RGB image
    age = Image.open(self.images[index]).convert('RGB')

    next load the target
    rget = Image.open(self.targets[index]).convert('L')

    convert to pytorch tensors
    target = TF.to_tensor(target)
    rget = torch.from_numpy(np.array(target, dtype=np.uint8))
    age = TF.to_tensor(image)

    convert the labels into a mask
    rgetrgb = self.mask_to_rgb(target)
    rgetmask = self.mask_to_class(target)
    rgetmask = targetmask.long()
    rgetrgb = targetrgb.long()

    finally return the image pair
    return image, targetmask, targetrgb
turn image, targetmask

```

In []:

```

cityscapes_path = "data/"

print('Loading Cityscapes val datasets...')
dataset = CityscapesDataset(cityscapes_path, split='val', mode='fine', augment=False)

print(dataset)

val_dataloader = torch.utils.data.DataLoader(dataset, batch_size=1, shuffle=True, num_workers=0)

```

```

Loading Cityscapes val datasets...
data/leftImg8bit/val
data/gtFine/val
Dataset CityscapesDataset
  Number of images: 500
  Split: val
  Mode: gtFine
  Augment: False
  Root Location: data/

```

In []:

```

images, targetmask = next(iter(val_dataloader))
print(images.shape)
print(targetmask.shape)

images = [ img.cuda() for img in images ]
print(type(images))
print(images)
predictions = model(images)

print('Prediction keys:', predictions[0].keys())
print('Boxes shape:', predictions[0]['boxes'].shape)
print('Labels shape:', predictions[0]['labels'].shape)
print('Scores shape:', predictions[0]['scores'].shape)
print('Masks shape:', predictions[0]['masks'].shape)

torch.Size([1, 3, 1024, 2048])
torch.Size([1, 1024, 2048])
<class 'list'>
[tensor([[[[0.3804, 0.3686, 0.3686, ..., 0.2745, 0.2784, 0.2745],
          [0.3608, 0.3608, 0.3608, ..., 0.2784, 0.2745, 0.2706],
          [0.3412, 0.3490, 0.3490, ..., 0.2784, 0.2706, 0.2627],
          ...,
          [0.2039, 0.1961, 0.1922, ..., 0.1608, 0.1608, 0.1569],
          [0.2039, 0.1961, 0.1922, ..., 0.1725, 0.1686, 0.1608],
          [0.2039, 0.1961, 0.1922, ..., 0.1843, 0.1843, 0.1804]],
          ...,
          [[0.3843, 0.3843, 0.3922, ..., 0.2902, 0.2863, 0.2824],
          [0.3922, 0.4039, 0.4039, ..., 0.2902, 0.2863, 0.2824],
          [0.4000, 0.4118, 0.4157, ..., 0.2863, 0.2784, 0.2745],
          ...,
          [0.2588, 0.2588, 0.2588, ..., 0.2078, 0.2078, 0.2078],
          [0.2588, 0.2588, 0.2588, ..., 0.2196, 0.2157, 0.2078],
          [0.2588, 0.2588, 0.2588, ..., 0.2314, 0.2314, 0.2275]],
          ...,
          [[0.3255, 0.3255, 0.3137, ..., 0.2353, 0.2314, 0.2235],
          [0.3412, 0.3608, 0.3490, ..., 0.2314, 0.2275, 0.2196],
          [0.3569, 0.3725, 0.3765, ..., 0.2275, 0.2235, 0.2157],
          ...,
          [0.2275, 0.2235, 0.2275, ..., 0.1765, 0.1765, 0.1765],
          [0.2275, 0.2235, 0.2275, ..., 0.1882, 0.1843, 0.1725],
          [0.2275, 0.2235, 0.2275, ..., 0.2039, 0.2000, 0.1961]]],
          device='cuda:0')]
```

```

/usr/local/lib/python3.7/dist-packages/torch/functional.py:445: UserWarning: torch.meshgrid: in an upcoming release, it will be required to pass the indexing argument. (Triggered internally at ../aten/src/ATen/native/TensorShape.cpp:2157.)
```

```

    return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]
```

```

Prediction keys: dict_keys(['boxes', 'labels', 'scores', 'masks'])
Boxes shape: torch.Size([56, 4])
Labels shape: torch.Size([56])
Scores shape: torch.Size([56])
Masks shape: torch.Size([56, 1, 1024, 2048])
```

In []:

```

import numpy as np
import cv2
import random

# Array of labels for COCO dataset (91 elements)

coco_names = [
    '__background__', 'person', 'bicycle', 'car', 'motorcycle', 'airplane', 'bus',
    'train', 'truck', 'boat', 'traffic light', 'fire hydrant', 'N/A', 'stop sign',
    'parking meter', 'bench', 'bird', 'cat', 'dog', 'horse', 'sheep', 'cow',
    'elephant', 'bear', 'zebra', 'giraffe', 'N/A', 'backpack', 'umbrella', 'N/A', 'N/A',
    'handbag', 'tie', 'suitcase', 'frisbee', 'skis', 'snowboard', 'sports ball',
    'kite', 'baseball bat', 'baseball glove', 'skateboard', 'surfboard', 'tennis racket',
    'bottle', 'N/A', 'wine glass', 'cup', 'fork', 'knife', 'spoon', 'bowl',
    'banana', 'apple', 'sandwich', 'orange', 'broccoli', 'carrot', 'hot dog', 'pizza',
    'donut', 'cake', 'chair', 'couch', 'potted plant', 'bed', 'N/A', 'dining table',
    'N/A', 'N/A', 'toilet', 'N/A', 'tv', 'laptop', 'mouse', 'remote', 'keyboard', 'cell phone',
    'microwave', 'oven', 'toaster', 'sink', 'refrigerator', 'N/A', 'book',
    'clock', 'vase', 'scissors', 'teddy bear', 'hair drier', 'toothbrush'
]

# Random colors to use for labeling objects

COLORS = np.random.uniform(0, 255, size=(len(coco_names), 3)).astype(np.uint8)

# Overlay masks, bounding boxes, and labels on input numpy image

def draw_segmentation_map(image, masks, boxes, labels):
    alpha = 1
    beta = 0.5 # transparency for the segmentation map
    gamma = 0 # scalar added to each sum
    # convert from RGB to OpenCV BGR format
    image = cv2.cvtColor(image, cv2.COLOR_RGB2BGR)
    for i in range(len(masks)):
        mask = masks[i,:,:]
        red_map = np.zeros_like(mask).astype(np.uint8)
        green_map = np.zeros_like(mask).astype(np.uint8)
        blue_map = np.zeros_like(mask).astype(np.uint8)
        # apply a random color mask to each object
        color = COLORS[random.randrange(0, len(COLORS))]
        red_map[mask > 0.5] = color[0]
        green_map[mask > 0.5] = color[1]
        blue_map[mask > 0.5] = color[2]
        # combine all the masks into a single image
        segmentation_map = np.stack([red_map, green_map, blue_map], axis=2)
        # apply colored mask to the image
        image = cv2.addWeighted(image, alpha, segmentation_map, beta, gamma)
        # draw the bounding box around each object
        p1 = (int(boxes[i][0]), int(boxes[i][1]))
        p2 = (int(boxes[i][2]), int(boxes[i][3]))
        color = (int(color[0]), int(color[1]), int(color[2]))
        cv2.rectangle(image, p1, p2, color, 2)
        # put the label text above the objects
        p = (int(boxes[i][0]), int(boxes[i][1]-10))
        cv2.putText(image, labels[i], p, cv2.FONT_HERSHEY_SIMPLEX, 0.5, color, 2, cv2.LINE_AA)

    return cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

# Overlay masks, bounding boxes, and labels of objects with scores greater than

```

```
# threshold on one of the images in the input tensor using the predictions output by
```

```
def prediction_to_mask_image(images, predictions, img_index, threshold):
    scores = predictions[img_index]['scores']
    boxes_to_use = scores >= threshold
    img = (images[img_index].cpu().permute(1, 2, 0).numpy() * 255).astype(np.uint8)
    masks = predictions[img_index]['masks'][boxes_to_use, :, :].cpu().detach().squeeze()
    boxes = predictions[img_index]['boxes'][boxes_to_use, :].cpu().detach().numpy()
    labels = predictions[img_index]['labels'][boxes_to_use].cpu().numpy()
    labels = [ coco_names[l] for l in labels ]

    return draw_segmentation_map(img, masks, boxes, labels)
```

```
In [ ]:
```

```
from matplotlib import pyplot as plt
```

```
masked_img = prediction_to_mask_image(images, predictions, 0, 0.5)
plt.figure(1, figsize=(12, 9), dpi=100)
plt.imshow(masked_img)
plt.title('Validation image result')
plt.show()
```



```
In [ ]:
```

coco-finetune

In []:

```

es_path = "data/"

loading Cityscapes val datasets...'
taset = CityscapesDataset(cityscapes_path, split='train', mode='fine', augment=False)
set = CityscapesDataset(cityscapes_path, split='val', mode='fine', augment=False)

rain dataset')
ain_dataset)
al dataset')
l_dataset)

dataset = get_cityscapes()

llate_fn(batch):
turn tuple(zip(*batch))

taloader = torch.utils.data.DataLoader(train_dataset, batch_size=1, shuffle=True, num
loader = torch.utils.data.DataLoader(val_dataset, batch_size=1, shuffle=True, num_wo:

```

Loading Cityscapes val datasets...

```

data/leftImg8bit/train
data/gtFine/train
data/leftImg8bit/val
data/gtFine/val
train dataset
Dataset CityscapesDataset
  Number of images: 2975
  Split: train
  Mode: gtFine
  Augment: False
  Root Location: data/

```

val dataset

```

Dataset CityscapesDataset
  Number of images: 500
  Split: val
  Mode: gtFine
  Augment: False
  Root Location: data/

```

In []:

```
import torch
import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor
from PIL import Image

num_classes = 91

print('loading pretrained model')
model = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained=True)

# Modify model for the given number of classes

in_features = model.roi_heads.box_predictor.cls_score.in_features
num_classes = 8

model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)

in_features_mask = model.roi_heads.mask_predictor.conv5_mask.in_channels
hidden_layer = 256
model.roi_heads.mask_predictor = MaskRCNNPredictor(in_features_mask,
                                                    hidden_layer,
                                                    num_classes)

model.cuda()
```

loading pretrained model

In []:

```

import numpy as np
import cv2
import random
import matplotlib.pyplot as plt

# Array of labels for COCO dataset (91 elements)

coco_names = [
    '__background__', 'person', 'bicycle', 'car', 'motorcycle', 'airplane', 'bus',
    'train', 'truck', 'boat', 'traffic light', 'fire hydrant', 'N/A', 'stop sign',
    'parking meter', 'bench', 'bird', 'cat', 'dog', 'horse', 'sheep', 'cow',
    'elephant', 'bear', 'zebra', 'giraffe', 'N/A', 'backpack', 'umbrella', 'N/A', 'N/A',
    'handbag', 'tie', 'suitcase', 'frisbee', 'skis', 'snowboard', 'sports ball',
    'kite', 'baseball bat', 'baseball glove', 'skateboard', 'surfboard', 'tennis racket',
    'bottle', 'N/A', 'wine glass', 'cup', 'fork', 'knife', 'spoon', 'bowl',
    'banana', 'apple', 'sandwich', 'orange', 'broccoli', 'carrot', 'hot dog', 'pizza',
    'donut', 'cake', 'chair', 'couch', 'potted plant', 'bed', 'N/A', 'dining table',
    'N/A', 'N/A', 'toilet', 'N/A', 'tv', 'laptop', 'mouse', 'remote', 'keyboard', 'cell phone',
    'microwave', 'oven', 'toaster', 'sink', 'refrigerator', 'N/A', 'book',
    'clock', 'vase', 'scissors', 'teddy bear', 'hair drier', 'toothbrush'
]

# Random colors to use for labeling objects

COLORS = np.random.uniform(0, 255, size=(len(coco_names), 3)).astype(np.uint8)

# Overlay masks, bounding boxes, and labels on input numpy image

def draw_segmentation_map(image, masks, boxes, labels):
    alpha = 1
    beta = 0.5 # transparency for the segmentation map
    gamma = 0 # scalar added to each sum
    # convert from RGB to OpenCV BGR format
    image = cv2.cvtColor(image, cv2.COLOR_RGB2BGR)
    for i in range(len(masks)):
        mask = masks[i,:,:]
        red_map = np.zeros_like(mask).astype(np.uint8)
        green_map = np.zeros_like(mask).astype(np.uint8)
        blue_map = np.zeros_like(mask).astype(np.uint8)
        # apply a random color mask to each object
        color = COLORS[random.randrange(0, len(COLORS))]
        red_map[mask > 0.5] = color[0]
        green_map[mask > 0.5] = color[1]
        blue_map[mask > 0.5] = color[2]
        # combine all the masks into a single image
        segmentation_map = np.stack([red_map, green_map, blue_map], axis=2)
        # apply colored mask to the image
        image = cv2.addWeighted(image, alpha, segmentation_map, beta, gamma)
        # draw the bounding box around each object
        p1 = (int(boxes[i][0]), int(boxes[i][1]))
        p2 = (int(boxes[i][2]), int(boxes[i][3]))
        color = (int(color[0]), int(color[1]), int(color[2]))
        cv2.rectangle(image, p1, p2, color, 2)
        # put the label text above the objects
        p = (int(boxes[i][0]), int(boxes[i][1]-10))
        cv2.putText(image, labels[i], p, cv2.FONT_HERSHEY_SIMPLEX, 0.5, color, 2, cv2.LINE_AA)

    return cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

```

```
# Overlay masks, bounding boxes, and labels of objects with scores greater than
# threshold on one of the images in the input tensor using the predictions output by
```

```
def prediction_to_mask_image(images, predictions, img_index, threshold):
    scores = predictions[img_index]['scores']
    boxes_to_use = scores >= threshold
    img = (images[img_index].cpu().permute(1, 2, 0).numpy() * 255).astype(np.uint8)
    masks = predictions[img_index]['masks'][boxes_to_use, :, :].cpu().detach().squeeze()
    boxes = predictions[img_index]['boxes'][boxes_to_use, :].cpu().detach().numpy()
    labels = predictions[img_index]['labels'][boxes_to_use].cpu().numpy()
    labels = [ coco_names[l] for l in labels ]

    return draw_segmentation_map(img, masks, boxes, labels)
```

Inference

In []:

```
!wget https://www.cs.ait.ac.th/~mdailey/20201112_072342.jpg
```

```
--2022-03-03 14:24:40-- https://www.cs.ait.ac.th/~mdailey/20201112_07
2342.jpg (https://www.cs.ait.ac.th/~mdailey/20201112_072342.jpg)
Resolving www.cs.ait.ac.th (www.cs.ait.ac.th)... 192.41.170.42
Connecting to www.cs.ait.ac.th (www.cs.ait.ac.th)|192.41.170.42|:44
3... connected.
HTTP request sent, awaiting response... 200 OK
Length: 3320786 (3.2M) [image/jpeg]
Saving to: '20201112_072342.jpg.1'
```

```
20201112_072342.jpg 100%[=====>] 3.17M 867KB/s in
3.7s
```

```
2022-03-03 14:24:45 (867 KB/s) - '20201112_072342.jpg.1' saved [332078
6/3320786]
```


In []:

```

#model = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained = True)
model_test = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained=True)
model_test.eval()

im = Image.open('20201112_072342.jpg').rotate(180)

print(type(images))
#images = [ img.cuda() for img in images ]
transform = T.ToTensor()
image_batch = transform(im).unsqueeze(0)
# #image_batch.shap
print(image_batch)
predictions = model_test(image_batch)

print('Prediction keys:', predictions[0].keys())
print('Boxes shape:', predictions[0]['boxes'].shape)
print('Labels shape:', predictions[0]['labels'].shape)
print('Scores shape:', predictions[0]['scores'].shape)
print('Masks shape:', predictions[0]['masks'].shape)

import numpy as np
import cv2
import random

# Array of labels for COCO dataset (91 elements)

coco_names = [
    '__background__', 'person', 'bicycle', 'car', 'motorcycle', 'airplane', 'bus',
    'train', 'truck', 'boat', 'traffic light', 'fire hydrant', 'N/A', 'stop sign',
    'parking meter', 'bench', 'bird', 'cat', 'dog', 'horse', 'sheep', 'cow',
    'elephant', 'bear', 'zebra', 'giraffe', 'N/A', 'backpack', 'umbrella', 'N/A', 'N/A',
    'handbag', 'tie', 'suitcase', 'frisbee', 'skis', 'snowboard', 'sports ball',
    'kite', 'baseball bat', 'baseball glove', 'skateboard', 'surfboard', 'tennis racket',
    'bottle', 'N/A', 'wine glass', 'cup', 'fork', 'knife', 'spoon', 'bowl',
    'banana', 'apple', 'sandwich', 'orange', 'broccoli', 'carrot', 'hot dog', 'pizza',
    'donut', 'cake', 'chair', 'couch', 'potted plant', 'bed', 'N/A', 'dining table',
    'N/A', 'N/A', 'toilet', 'N/A', 'tv', 'laptop', 'mouse', 'remote', 'keyboard', 'cell phone',
    'microwave', 'oven', 'toaster', 'sink', 'refrigerator', 'N/A', 'book',
    'clock', 'vase', 'scissors', 'teddy bear', 'hair drier', 'toothbrush'
]

# Random colors to use for labeling objects

COLORS = np.random.uniform(0, 255, size=(len(coco_names), 3)).astype(np.uint8)

# Overlay masks, bounding boxes, and labels on input numpy image

def draw_segmentation_map(image, masks, boxes, labels):
    alpha = 1
    beta = 0.5 # transparency for the segmentation map
    gamma = 0 # scalar added to each sum
    # convert from RGB to OpenCV BGR format
    image = cv2.cvtColor(image, cv2.COLOR_RGB2BGR)
    for i in range(len(masks)):
        mask = masks[i,:,:]
        red_map = np.zeros_like(mask).astype(np.uint8)

```

```

green_map = np.zeros_like(mask).astype(np.uint8)
blue_map = np.zeros_like(mask).astype(np.uint8)
# apply a random color mask to each object
color = COLORS[random.randrange(0, len(COLORS))]
red_map[mask > 0.5] = color[0]
green_map[mask > 0.5] = color[1]
blue_map[mask > 0.5] = color[2]
# combine all the masks into a single image
segmentation_map = np.stack([red_map, green_map, blue_map], axis=2)
# apply colored mask to the image
image = cv2.addWeighted(image, alpha, segmentation_map, beta, gamma)
# draw the bounding box around each object
p1 = (int(boxes[i][0]), int(boxes[i][1]))
p2 = (int(boxes[i][2]), int(boxes[i][3]))
color = (int(color[0]), int(color[1]), int(color[2]))
cv2.rectangle(image, p1, p2, color, 2)
# put the label text above the objects
p = (int(boxes[i][0]), int(boxes[i][1]-10))
cv2.putText(image, labels[i], p, cv2.FONT_HERSHEY_SIMPLEX, 0.5, color, 2, cv

return cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

# Overlay masks, bounding boxes, and labels of objects with scores greater than
# threshold on one of the images in the input tensor using the predictions output by

def prediction_to_mask_image(images, predictions, img_index, threshold):
    scores = predictions[img_index]['scores']
    boxes_to_use = scores >= threshold
    img = (images[img_index].cpu().permute(1, 2, 0).numpy() * 255).astype(np.uint8)
    masks = predictions[img_index]['masks'][boxes_to_use, :, :].cpu().detach().squeeze()
    boxes = predictions[img_index]['boxes'][boxes_to_use, :].cpu().detach().numpy()
    labels = predictions[img_index]['labels'][boxes_to_use].cpu().numpy()
    labels = [ coco_names[l] for l in labels ]

    return draw_segmentation_map(img, masks, boxes, labels)

from matplotlib import pyplot as plt

masked_img = prediction_to_mask_image(image_batch, predictions, 0, 0.5)
plt.figure(1, figsize=(12, 9), dpi=100)
plt.imshow(masked_img)
plt.title('Validation image result')
plt.show()

```

```

<class 'list'>
tensor([[[[0.8314, 0.8314, 0.8353, ..., 0.8275, 0.8275, 0.8314],
          [0.8275, 0.8314, 0.8314, ..., 0.8196, 0.8196, 0.8235],
          [0.8314, 0.8353, 0.8353, ..., 0.8314, 0.8235, 0.8235],
          ...,
          [0.3216, 0.3137, 0.2902, ..., 0.1961, 0.1843, 0.1451],
          [0.2902, 0.3098, 0.2745, ..., 0.1255, 0.1098, 0.1255],
          [0.2667, 0.2784, 0.2784, ..., 0.1176, 0.1176, 0.2314]],
        [[0.8510, 0.8510, 0.8549, ..., 0.8471, 0.8471, 0.8510],
          [0.8471, 0.8510, 0.8510, ..., 0.8392, 0.8392, 0.8431],
          [0.8510, 0.8549, 0.8549, ..., 0.8510, 0.8431, 0.8431],
          ...,
          [0.3608, 0.3529, 0.3294, ..., 0.2118, 0.2000, 0.1608],
          [0.3294, 0.3490, 0.3137, ..., 0.1412, 0.1255, 0.1412]],

```

```

[0.3059, 0.3176, 0.3176, ..., 0.1333, 0.1333, 0.2471]],

[[0.8745, 0.8745, 0.8784, ..., 0.8706, 0.8706, 0.8745],
 [0.8706, 0.8745, 0.8745, ..., 0.8627, 0.8627, 0.8667],
 [0.8745, 0.8784, 0.8784, ..., 0.8745, 0.8667, 0.8667],
 ...,
 [0.3255, 0.3176, 0.2941, ..., 0.2235, 0.2118, 0.1725],
 [0.2941, 0.3137, 0.2784, ..., 0.1529, 0.1373, 0.1529],
 [0.2706, 0.2824, 0.2824, ..., 0.1451, 0.1451, 0.2588]]])
Prediction keys: dict_keys(['boxes', 'labels', 'scores', 'masks'])
Boxes shape: torch.Size([34, 4])
Labels shape: torch.Size([34])
Scores shape: torch.Size([34])
Masks shape: torch.Size([34, 1, 3024, 4032])

```



In []: